
CS4070 - MULTIVARIATE DATA ANALYSIS - PART 2

ASSIGNMENT 2

Konstantinos Krachtopoulos

Student Number: 5472539

Faculty of Electrical Engineering, Mathematics and Computer Science (EEMCS)

Technical University Delft

December 2021

1 Exercise 1

The log likelihood of the regression model, assuming that all y_i are independent is:

$$\begin{aligned}
 L(\theta; y) &= \prod_{i=1}^n e^{-\mu_i} \mu_i^{k_i} / (k_i!) \Rightarrow \\
 \log(L(\theta; y)) &= \log\left(\prod_{i=1}^n e^{-\mu_i} \mu_i^{k_i} / (k_i!)\right) \Rightarrow \\
 \log(L(\theta; y)) &= \sum_{i=1}^n (\log(e^{-\mu_i}) + \log(\mu_i^{k_i}) - \log(k_i!)) \Rightarrow \\
 \log(L(\theta; y)) &= \sum_{i=1}^n (-\mu_i + k_i \log(\mu_i) - \log(k_i!)) \Rightarrow
 \end{aligned} \tag{1}$$

Since $\mu_i = e^{x_i^T \theta}$, and $k = y_i$ the log likelihood becomes:

$$\log(L(\theta; y)) = \sum_{i=1}^n (-e^{x_i^T \theta} + y_i x_i^T \theta - \log(y_i!)) \tag{2}$$

2 Exercise 2

The log likelihood gradient is the first partial derivative of log likelihood with respect to θ :

$$\frac{\partial}{\partial \theta} \log(L(\theta; y)) = \sum_{i=1}^n (-x_i^T e^{x_i^T \theta} + y_i x_i^T) \tag{3}$$

The Hessian of the log likelihood is the second partial derivative with respect to θ and θ^T :

$$\frac{\partial^2}{\partial \theta \partial \theta^T} \log(L(\theta; y)) = \sum_{i=1}^n (-x_i^T x_i e^{x_i^T \theta}) \tag{4}$$

3 Exercise 3

Based on the Newton Algorithm, can be updated iteratively using the following formula:

$$\theta^{j+1} = \theta^j - H_L(\theta^j)^{-1} \nabla L(\theta^j) \tag{5}$$

The Hessian Matrix and the Gradient of the Log Likelihood have already been calculated in Exercise 2. The initial θ will be sampled from a normal distribution $N \sim (0, \tilde{\sigma}^2 I_p)$, where $\tilde{\sigma} = 4$.

In order to compute the posterior distribution of θ , and because we know that its distribution is normal, we can use Laplace approximation. With Laplace Approximation, the posterior distribution is calculated as

$$f(\theta|x)_{\theta|x} \approx \Phi(\theta, \tilde{\Theta}, -H_L(\tilde{\Theta})) \tag{6}$$

where $\tilde{\Theta}$ is the posterior mode of θ , calculated with Newton Algorithm.

After implementing the above formulas in Python, the following values are derived for the mean vector and the covariance matrix:

$$m = \begin{bmatrix} 1.12777336 \\ 0.42858431 \\ 0.01512131 \\ -0.05416601 \end{bmatrix}, \quad (7)$$

$$Cov = \begin{bmatrix} 0.03140956 & -0.00820031 & 0.00116572 & -0.00139328 \\ -0.00820031 & 0.00305825 & -0.00031134 & 0.00066138 \\ 0.00116572 & -0.00031134 & 0.0148073 & -0.0014676 \\ -0.00139328 & 0.00066138 & -0.0014676 & 0.01173136 \end{bmatrix} \quad (8)$$

4 Exercise 4

In order to implement Random Walk Metropolis-Hastings algorithm, we need to iteratively update θ with

$$\theta^o := \theta + \sigma_{proposal} N(0, I_p) \quad (9)$$

where the initial θ is sampled from the distribution $N(0, 4^2 I_p)$. The update takes place based on the Metropolis-Hastings acceptance probability:

$$\alpha(\theta, \theta^0) = \min(1, \frac{\pi(\theta^0) q(\theta^0, \theta)}{\pi(\theta) q(\theta, \theta^0)}) \quad (10)$$

where $q(., .)$ is the transition probability between two states, and $\pi(\theta)$ is the posterior probability of θ . Because we consider proposals of θ coming from a normal distribution, the transition matrix is symmetric, i.e. $q(\theta, \theta^0) = q(\theta^0, \theta)$. Hence, the previous equation becomes:

$$\alpha(\theta, \theta^0) = \min(1, \frac{\pi(\theta^0)}{\pi(\theta)}) \quad (11)$$

Moreover, in order to decide whether θ should be updated in the current iteration, we follow the relation:

$$\theta_{n+1} = \begin{cases} \theta^0, & \text{with prob } \alpha(\theta, \theta_0) \\ \theta, & \text{with prob } 1 - \alpha(\theta, \theta_0) \end{cases} \quad (12)$$

To implement the above equations, we calculate the log of the acceptance as $\log A = \log \pi(\theta^0) - \log \pi(\theta)$, and then compare it with u , where u is a sampled from $Unif(0, 1)$.

1. If $\log A > \log u$, the acceptance criterion isn't satisfied.
2. Else, the acceptance criterion isn't satisfied.

Finally, the parameter $\sigma_{proposal}$ needs tuning, in order to get an acceptance rate of 25-50%. After trying multiple values, the final value of $\sigma_{proposal}$ is set to be 0.08, resulting in an acceptance rate of 32.8%. The following graph shows the values of θ_2 versus θ_1 , where the color indicates the iteration number. Green indicates the first iterations, while blue indicates the last iterations.

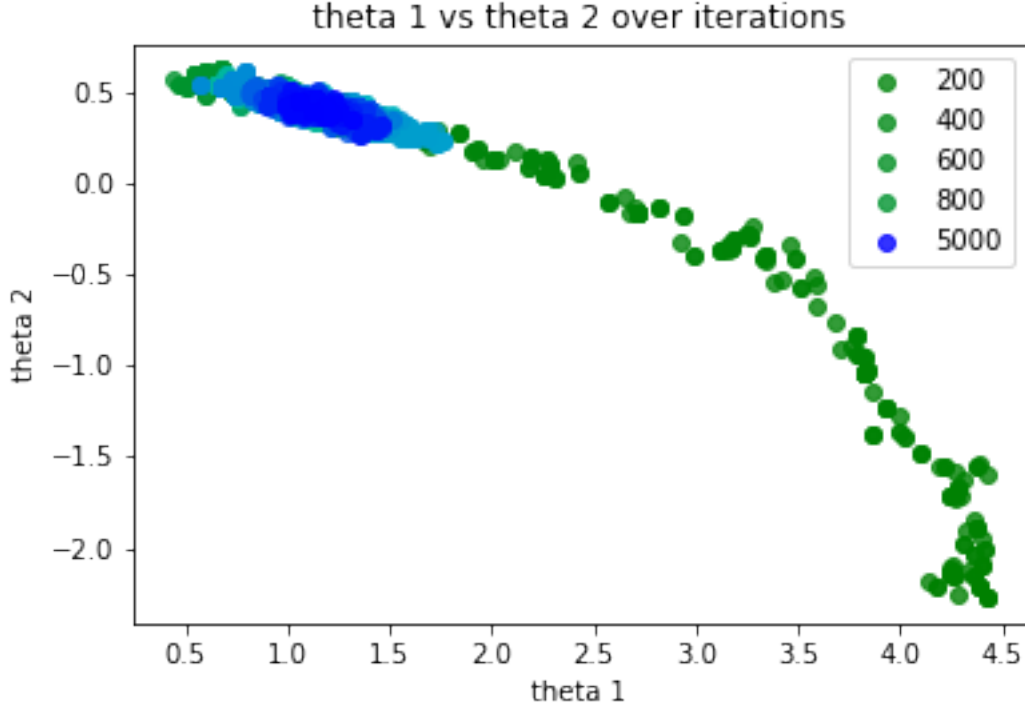


Figure 1: Random-walk iterations, θ_1, θ_2 comparison.

By checking the values of θ over the 5000 iterations, we see that its values started converging after 280 iterations. Hence, "burning" these iterations, the Monte Carlo posterior mean is:

$$\theta_{mean} = \begin{bmatrix} 1.1471 \\ 0.4196 \\ 0.0197 \\ -0.0376 \end{bmatrix} \quad (13)$$

5 Exercise 5

We want to iteratively sample from the conditionals of θ and $\tilde{\sigma}^2$. The sampling for θ has already been implemented in the previous exercise. Moreover, since $\tilde{\sigma}^2$ follows an inverse Gamma distribution with, its density function is given by:

$$p(\tilde{\sigma}^2 | \alpha, \beta) = \frac{\beta^\alpha}{\Gamma(\alpha)} (\tilde{\sigma}^2)^{-\alpha-1} \exp\left(-\frac{\beta}{\tilde{\sigma}^2}\right) \quad (14)$$

For the conditional distribution $\tilde{\sigma}^2 | \theta, y$:

$$\begin{aligned} p(\tilde{\sigma}^2 | \theta, y) &\propto p(y, \theta, \tilde{\sigma}^2) = \\ &(\tilde{\sigma}^2)^{-p/2} \exp\left(-\frac{1}{2\tilde{\sigma}^2} \|\theta\|^2\right) (\tilde{\sigma}^2)^{-A-1} \exp\left(-\frac{\beta}{\tilde{\sigma}^2}\right) \mathbf{1}_{(0, \infty)}(\tilde{\sigma}^2) \Rightarrow \\ &\tilde{\sigma}^2 | \theta, y \sim IG(\alpha + p/2, \beta + \|\theta\|^2 / 2) \end{aligned} \quad (15)$$

Hence, we follow the same approach as in Exercise 4. The only difference is that before sampling for θ , we sample for $\tilde{\sigma}^2$. The Traceplot below shows the values of $\tilde{\sigma}^2$ over the 5000 iterations.

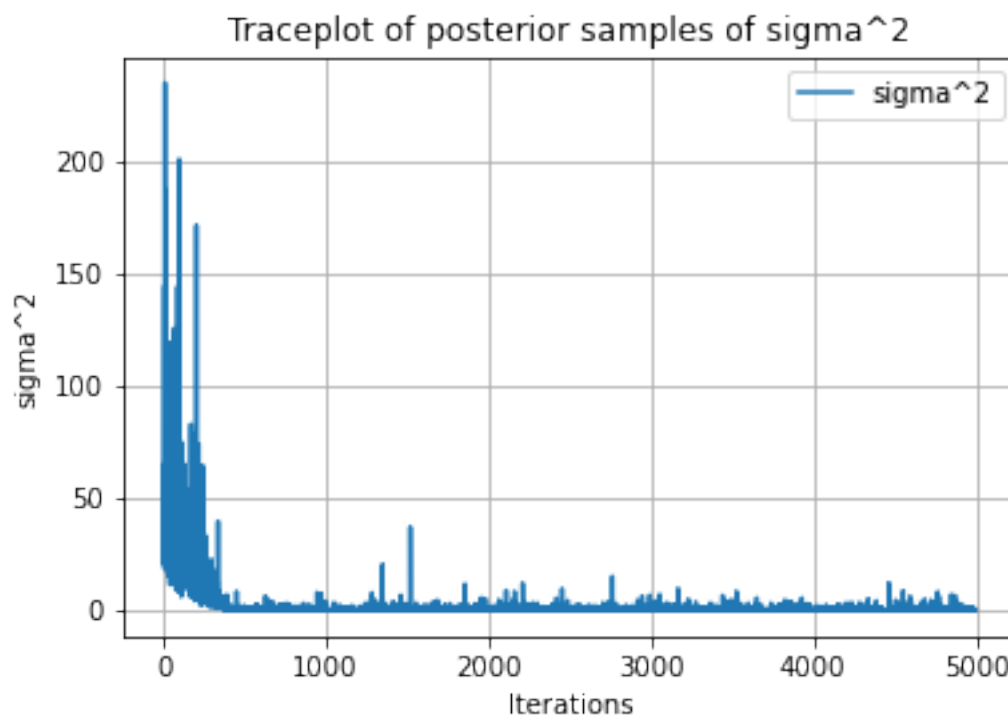


Figure 2: Traceplot of posterior samples of $\tilde{\sigma}^2$.

6 Appendix

The code used for computing the quantities of Exercises 3-4-5 is presented below, together with the appropriate comments:

6.1 Code - Exercise 3

```
# basic imports
import math

import numpy as np
import pandas as pd
from numpy.linalg import inv
import matplotlib.pyplot as plt

# Define functions to calculate fundamental values
def calc_gradient(x, y, thetas):
    # Calculate the gradient of the log likelihood
    return -x.T @ np.exp(x @ thetas.T) + x.T @ y

def calc_hessian_inverse(x, thetas):
    # Calculate the hessian matrix of the log likelihood
    hess = -x.T @ np.diag(np.exp(x @ thetas.T)) @ x
    return np.linalg.inv(hess)

# Load the data as a pandas dataframe
```

```

df = pd.read_csv("dataexercise2.csv")

# Convert the dataframe to a numpy array (for easier manipulation)
arr = df.to_numpy()

# Split the dataset
# first 4 columns are features (x matrix)
# last column is label (y vector)
x = arr[:, :-1]
y = arr[:, -1]

# Sample thetas from prior distribution
featNum = x.shape[1]
std0 = 4
thetas = np.random.multivariate_normal(np.zeros(featNum), std0**2*np.eye(featNum))

# Newton algorithm implementation
iters = 1000
e = 1e-6

for i in range(iters):
    grad = calc_gradient(x, y, thetas)
    hess_inv = calc_hessian_inverse(x, thetas)
    update = grad @ hess_inv
    if np.any(np.abs(update)) > e:
        thetas = thetas - update
    else:
        print("Converged after {} iterations".format(i + 1))
        break

# Calculate the gaussian distribution of thetas with Laplace approximation
mean = thetas
cov = -calc_hessian_inverse(x, thetas)
print(mean)
print(cov)

```

6.2 Code - Exercise 4

```

# basic imports
import numpy as np
import pandas as pd
import math
import matplotlib.pyplot as plt
from scipy.stats import multivariate_normal
from colour import Color

# Define functions to modularize the script
def calcPriorTheta(thetas, std0=4, featNum=4):
    # Calculate the prior of thetas, based on normal distribution
    # thetas: 4x1 array of parameters
    # std0: standard deviation set by the problem statement
    # featNum: Number of features loaded
    dist = multivariate_normal(mean=np.zeros(featNum), cov=std0**2*np.eye(featNum))
    return dist.pdf(thetas.T)

```

```

def calcLogPostTheta(thetas, x, y):
    # Calculate the posterior of theta given data, without the
    # regularization parameter
    # thetas: 4x1 array of parameters
    # x: 2d given dataset
    # y: 1d labels of given dataset
    log_likelihood = 0
    for i in range(y.shape[0]):
        log_likelihood += -np.exp(x[i].T @ thetas)\
            + y[i]*x[i].T @ thetas \
            - np.log(math.factorial(y[i]))
    log_prior = calcPriorTheta(thetas)
    return (log_likelihood + log_prior)[0]

def acceptUpdate(logPostThetasCur, logPostThetasNew):
    # Calculate whether currents thetas will be updated with the new thetas
    # based on the acceptance criterion.
    # logPostThetasCur
    # logPostThetasNew
    logAcceptance = logPostThetasNew - logPostThetasCur
    if logAcceptance > 0: # then a = 1, sure update
        return True
    else:
        # Sample from uniform distribution
        logU = np.log(np.random.uniform(0, 1))
        return logAcceptance > logU

def execMetropolisHastings(thetasInit, iters, x, y):
    # logPosterior: the log posterior probability of thetas
    # thetasInit: initialization of theta parameters
    # iters: number of iterations
    # x: 2d given dataset
    # y: 1d labels of given dataset

    # Create lists to store the accepted, rejected and current thetas
    acc_thetas = []
    rej_thetas = []
    cur_thetas = np.zeros((iters, featNum))

    # Initialize theta parameters
    thetasCur = np.reshape(thetasInit, (-1, 1))

    # Initialized proposed sigma
    sigma_prop = 0.08

    # Begin the loop
    for iter in range(iters):
        # Compute the new theta parameters with Random Walk
        thetasNew = thetasCur + np.random.normal(0, sigma_prop, (featNum, 1))
        thetasCurPost = calcLogPostTheta(thetasCur, x, y)
        thetasNewPost = calcLogPostTheta(thetasNew, x, y)
        if acceptUpdate(thetasCurPost, thetasNewPost):
            # Update thetas

```

```

        thetasCur = thetasNew
        acc_thetas.append(thetasNew)
    else:
        rej_thetas.append(thetasNew)
    # Collect current theta parameters
    cur_thetas[iter] = thetasCur[:, 0]

    return cur_thetas, acc_thetas, rej_thetas

# Main code execution
# Load the data as a pandas dataframe
df = pd.read_csv("dataexercise2.csv")

# Convert the dataframe to a numpy array (for easier manipulation)
arr = df.to_numpy()

# Split the dataset
# first 4 columns are features (x matrix)
# last column is label (y vector)
x = arr[:, :-1]
y = arr[:, -1]

# Set initial parameters
featNum = x.shape[1]
std0 = 4
thetasInit = np.random.multivariate_normal(np.zeros(featNum), std0**2*np.eye(featNum))
iters = 5000

cur_thetas, acc_thetas, rej_thetas = \
    execMetropolisHastings(thetasInit, iters, x, y)

acceptance_rate = len(acc_thetas)/iters
print("Acceptance rate is {}".format(acceptance_rate))

# plot theta1 vs theta2
initCol = Color("green")
finCol = Color("blue")
colorSpectrum = list(initCol.range_to(finCol, len(cur_thetas)))
fig, ax = plt.subplots()
for i in range(len(colorSpectrum)):
    if (((i+1) % 200 == 0) and (i+1 < 1000)) or (i+1 == iters):
        ax.scatter(cur_thetas[i, 0], cur_thetas[i, 1], alpha=0.8, c=str(colorSpectrum[i]),
                    label=i+1)
    else:
        ax.scatter(cur_thetas[i, 0], cur_thetas[i, 1], alpha=0.8, c=str(colorSpectrum[i]))

plt.xlabel("theta 1")
plt.ylabel("theta 2")
plt.title("theta 1 vs theta 2 over iterations")
plt.legend()
plt.show()
fig.savefig('plot4.png')

# Compute the mean posterior

```



```

lim = 280
meanPost = np.mean(cur_thetas[280:, :], axis=0)
print("Mean posterior theta is: {}".format(meanPost))

```

6.3 Code - Exercise 5

```

# basic imports
import numpy as np
import pandas as pd
import math
import matplotlib.pyplot as plt
from scipy.stats import multivariate_normal
from scipy.stats import invgamma

# Define functions to modularize the script
def calcPriorTheta(thetas, std0=4, featNum=4):
    # Calculate the prior of thetas, based on normal distribution
    # thetas: 4x1 array of parameters
    # std0: standard deviation set by the problem statement
    # featNum: Number of features loaded
    dist = multivariate_normal(mean=np.zeros(featNum), cov=std0**2*np.eye(featNum))
    return dist.pdf(thetas.T)

def calcLogPostTheta(thetas, x, y, sigma=4):
    # Calculate the posterior of theta given data, without the
    # regularization parameter
    # thetas: 4x1 array of parameters
    # x: 2d given dataset
    # y: 1d labels of given dataset
    log_likelihood = 0
    for i in range(y.shape[0]):
        log_likelihood += -np.exp(x[i].T @ thetas)\
            + y[i]*x[i].T @ thetas\
            - np.log(math.factorial(y[i]))
    log_prior = calcPriorTheta(thetas, std0=sigma)
    return (log_likelihood + log_prior)[0]

def acceptUpdate(logPostThetasCur, logPostThetasNew):
    # Calculate whether currents thetas will be updated with the new thetas
    # based on the acceptance criterion.
    # logPostThetasCur
    # logPostThetasNew
    logAcceptance = logPostThetasNew - logPostThetasCur
    if logAcceptance > 0: # then a = 1, sure update
        return True
    else:
        # Sample from uniform distribution
        logU = np.log(np.random.uniform(0, 1))
        return logAcceptance > logU

def execGibbsSamples(thetasInit, iters, x, y):
    # logPosterior: the log posterior probability of thetas
    # thetasInit: initialization of theta parameters
    # iters: number of iterations

```

```

    # x: 2d given dataset
    # y: 1d labels of given dataset

    # Create lists to store the accepted, rejected and current thetas
    acc_thetas = []
    rej_thetas = []
    cur_thetas = np.zeros((iters, featNum))
    sigmas = []

    # Initialize theta parameters
    thetasCur = np.reshape(thetasInit, (-1, 1))

    # Initialized proposed sigma
    sigma_prop = 0.08
    alpha = 0.2
    beta = 0.2

    # Begin the loop
    for iter in range(iters):
        # First update the sigma, based on the inversed gamma distribution
        sigma = invgamma.rvs(a=alpha + featNum/2, scale=beta + 0.5*np.linalg.norm(thetasCur)**2)
        sigmas.append(sigma)
        # Compute the new theta parameters with Random Walk
        thetasNew = thetasCur + np.random.normal(0, sigma_prop, (featNum, 1))
        thetasCurPost = calcLogPostTheta(thetasCur, x, y, sigma=sigma)
        thetasNewPost = calcLogPostTheta(thetasNew, x, y, sigma=sigma)
        if acceptUpdate(thetasCurPost, thetasNewPost):
            # Update thetas
            thetasCur = thetasNew
            acc_thetas.append(thetasNew)
        else:
            rej_thetas.append(thetasNew)
            # Collect current theta parameters
            cur_thetas[iter] = thetasCur[:, 0]

    return cur_thetas, acc_thetas, rej_thetas, sigmas

# Main code execution
# Load the data as a pandas dataframe
df = pd.read_csv("dataexercise2.csv")

# Convert the dataframe to a numpy array (for easier manipulation)
arr = df.to_numpy()

# Split the dataset
# first 4 columns are features (x matrix)
# last column is label (y vector)
x = arr[:, :-1]
y = arr[:, -1]

# Set initial parameters
featNum = x.shape[1]
std0 = 4
thetasInit = np.random.multivariate_normal(np.zeros(featNum), std0**2*np.eye(featNum))

```

```
iters = 5000

cur_thetas, acc_thetas, rej_thetas, sigmas = \
    execGibbsSamples(thetasInit, iters, x, y)

acceptance_rate = len(acc_thetas)/iters
print("Acceptance rate is {}".format(acceptance_rate))

# plot traceplot
plt.plot(np.arange(len(sigmas)), sigmas, label="sigma^2")
plt.xlabel("Iterations")
plt.ylabel("sigma^2")
plt.title("Traceplot of posterior samples of sigma^2")
plt.legend()
plt.grid()
plt.savefig('plot5_1.png')
plt.show()
```